

Predicting User Sentiment in Software Issue Reports

ST2MLE — Machine Learning for IT Engineers · Project Report

1. Project overview

We build and evaluate a machine-learning pipeline that classifies the sentiment of user-reported software issues into **negative**, **neutral** and **positive**. The study compares **four text-vectorisation strategies** (Bag-of-Words, TF-IDF, corpus-trained Word2Vec, pre-trained GloVe — plus an optional BERT path) against **five classifiers** (Naive Bayes, Decision Tree, Random Forest, AdaBoost, Logistic Regression), each hyper-parameter-tuned with 5-fold cross-validation. We then study, in dedicated experiments, class-imbalance handling, dimensionality reduction, decision-tree pruning and early stopping.

All numbers in this report are produced by `python main.py` and are fully reproducible (single global seed `RANDOM_STATE = 42`).

2. Dataset description

2.1 Source

We use a **real, pre-cleaned review dataset** — the "App Store reviews" option of the brief. It is a Kaggle export of **Facebook app reviews** (10 000 reviews), each row being the review content and a **1-5 score**. The CSV ships with the repo ([data/facebook_reviews.csv](#)) and is loaded by [src/real_data.py](#) .

Sentiment labels are derived from the star score with the standard convention:

Score	Sentiment
1-2 ★	negative
3 ★	neutral
4-5 ★	positive


Two alternative sources are also wired in via `config.DATASET_SOURCE` : the Google Play `sealuzh/user_quality` corpus ("app_reviews" , 288 k reviews, downloaded on demand) and a reproducible **synthetic generator** ("synthetic" , offline fallback). The pipeline is identical for all three.

2.2 Composition (5 923 unique reviews)

We **de-duplicate** the review text — the same short review ("good", "nice") recurs thousands of times, and dropping duplicates prevents the identical string leaking across the train/test

split. After de-duplication and dropping reviews that clean to empty text, **5 923** remain, keeping the natural positive skew:

Class	Source score	Count	Share
negative	1-2 ★	1 857	31.4 %
neutral	3 ★	265	4.5 %
positive	4-5 ★	3 801	64.2 %



This real, severe imbalance (positives outnumber neutrals ~14:1) is exactly what the resampling experiment in §6 addresses.

2.3 Why the task is genuinely hard

Real Facebook reviews bring real difficulty — no synthetic noise needed:

1. **Star-derived labels are a proxy.** A 3★ "neutral" review is often mildly positive or negative in wording, so the tiny `neutral` class genuinely overlaps its neighbours.
 2. **Short, noisy, multilingual text.** Reviews are very short and contain typos, emoji and many non-English fragments (Spanish, Bengali...); the latter mostly clean to empty and are dropped.
 3. **Extreme imbalance.** With only 4.5 % `neutral`, a classifier can reach high accuracy while essentially never predicting the minority — making macro-F1 the metric that matters (see §5-§6).
-

3. ML pipeline overview

raw text → preprocess → vectorise → [balance] → [reduce] → classify → evaluate

3.1 Preprocessing ([src/preprocessing.py](#))

Lower-casing → URL/HTML/punctuation stripping → tokenisation → stop-word removal → WordNet lemmatisation. Two sentiment-aware choices:

- **Negations are kept** (`not` , `no` , `never` ...): they flip sentiment, so removing them as ordinary stop-words would destroy signal.
- **Two-pass lemmatisation** (verb then noun) normalises "crashing → crash", "issues → issue" cheaply, without a costly POS tagger.

3.2 Vectorisation ([src/vectorization.py](#))

Strategy	Representation	Notes
Bag-of-Words	n-gram counts (1-2-grams, ≤5 000 features)	sparse, non-negative
TF-IDF	counts re-weighted by inverse document frequency, sublinear TF	sparse, non-negative
Word2Vec	mean of 100-d embeddings trained on the training corpus	dense
GloVe (pre-trained)	mean of 100-d Wikipedia+Gigaword vectors (400 k vocab)	dense
BERT (optional)	mean-pooled contextual embeddings	requires <code>transformers + torch</code>

GloVe gives us a genuinely **pre-trained** embedding (loaded via `gensim-data`); the same vectoriser switches to the canonical 300-d Google-News word2vec vectors by changing one argument.

BERT is implemented in full (`BertVectorizer`) but needs the optional deep-learning stack and network access to the model hub. Where those are unavailable it is skipped automatically and the experiments use BoW / TF-IDF / Word2Vec / GloVe.

3.3 Classifiers ([src/models.py](#))

Naive Bayes (`MultinomialNB` for sparse, `GaussianNB` for dense — chosen automatically), Decision Tree (with pre- and post-pruning in the grid), Random Forest, AdaBoost, and Logistic Regression (used to demonstrate **L2 regularisation** via the tuned `C`).

3.4 Evaluation ([src/evaluation.py](#))

Accuracy, **macro**-averaged precision/recall/F1 (so minority classes count equally), and confusion matrices. Macro-F1 is the tuning target.

3.5 Two engineering choices that keep it fast and correct

- **Preprocess once.** Cleaning is stateless, so running it a single time up front is exactly equivalent to running it inside every CV fold — but far cheaper.
- **Cache vectorisers** (`joblib.Memory`) against one fixed `StratifiedKFold` , so the expensive Word2Vec / TF-IDF fits are computed once per fold and reused across every hyper-parameter combination and classifier. The full study (4 vectorisers × 5 classifiers + the four ablations) runs in ~8 min on 4 cores.

4. Experimental setup


- **Split:** stratified 80 % train / 20 % test (`TEST_SIZE = 0.20`).
- **Cross-validation:** 5-fold `StratifiedKFold` on the training set.
- **Tuning:** `GridSearchCV` , scoring `f1_macro` , refit on the best configuration.
- **Reproducibility:** every stochastic component seeded from `RANDOM_STATE = 42` ; Word2Vec uses a single worker for determinism.

5. Model comparison


Test-set results (held-out 20 %), sorted by macro-F1:

Rank	Vectoriser	Classifier	CV F1	Accuracy	Precision	Recall	F1 (macro)
1	TF-IDF	Naive Bayes	0.550	0.839	0.550	0.569	0.559
2	TF-IDF	Logistic Regression	0.557	0.830	0.547	0.563	0.555
3	TF-IDF	Random Forest	0.541	0.824	0.540	0.558	0.548
4	BoW	Naive Bayes	0.569	0.817	0.542	0.554	0.548
5	BoW	Logistic Regression	0.561	0.806	0.550	0.545	0.546
6	BoW	Decision Tree	0.533	0.768	0.547	0.533	0.539
7	BoW	Random Forest	0.542	0.810	0.534	0.542	0.537
8	Word2Vec	Logistic Regression	0.532	0.810	0.527	0.547	0.537
9	GloVe	Logistic Regression	0.537	0.800	0.521	0.543	0.532
10	Word2Vec	Random Forest	0.539	0.801	0.520	0.542	0.530
11	GloVe	Random Forest	0.531	0.805	0.532	0.531	0.528
12	TF-IDF	Decision Tree	0.527	0.784	0.578	0.530	0.528
13	Word2Vec	Naive Bayes	0.517	0.685	0.542	0.537	0.518
14	Word2Vec	AdaBoost	0.518	0.780	0.503	0.531	0.516
15	GloVe	AdaBoost	0.520	0.774	0.500	0.516	0.508
16	Word2Vec	Decision Tree	0.533	0.763	0.496	0.514	0.505
17	GloVe	Decision Tree	0.503	0.726	0.504	0.498	0.501
18	GloVe	Naive Bayes	0.459	0.678	0.504	0.518	0.497

Rank	Vectoriser	Classifier	CV F1	Accuracy	Precision	Recall	F1 (macro)
19	TF-IDF	AdaBoost	0.459	0.757	0.512	0.471	0.472
20	BoW	AdaBoost	0.454	0.753	0.516	0.463	0.465



Confusion matrix of the best model (TF-IDF + Naive Bayes):



Reading the results

First, a crucial caveat: **accuracy is misleading here**. Because ~64 % of reviews are positive and only 4.5 % are neutral, a model can score ~0.84 accuracy while **never predicting the neutral class at all** — so **macro-F1 (which weights all three classes equally) is the metric that matters**, and it sits around 0.46-0.56.


- **TF-IDF + Naive Bayes wins (macro-F1 = 0.559)**. On short, noisy review text, IDF-weighted counts plus NB's robustness to high dimensionality generalise best; BoW + NB is a near-tie (0.548).
- **The sparse count methods (TF-IDF/BoW) fill the top seven**, with Naive Bayes, Logistic Regression and Random Forest all close — linear/probabilistic models suit short lexical text, and the scores cluster tightly (0.54-0.56) because every model struggles equally with the tiny neutral class.
- **AdaBoost is worst (0.46-0.47)**: its depth-1 stumps can each test only one word out of thousands; the dense embeddings help it only marginally.
- **Embeddings (Word2Vec, GloVe) sit mid-table (0.50-0.54)**. Mean-pooled vectors wash out the specific lexical cues ("crash", "love", "buggy") short reviews live on; Gaussian NB on GloVe is near the bottom (0.497). A clear reminder that pre-trained is not automatically better than in-domain features.

- **Scores are far lower and tighter than on clean/synthetic data** — the honest signature of a real, noisy, severely imbalanced corpus with proxy (star-derived) labels.

6. Class-imbalance handling

Fixed pipeline (TF-IDF + Logistic Regression); only the resampling strategy changes. Resampling is applied **inside an imblearn pipeline**, i.e. to the training folds only.

Strategy	Accuracy	Macro-F1	Recall neg	Recall neu	Recall pos
none (baseline)	0.836	0.555	0.745	0.000	0.939
SMOTE	0.727	0.551	0.758	0.189	0.749
under-sampling	0.601	0.505	0.637	0.472	0.592



Interpretation — this is the most striking result in the study. The untreated baseline reaches a deceptively high **0.836 accuracy** while **never once predicting the neutral class** — its `neutral` recall is exactly **0.000** — because with only 4.5 % neutrals the model maximises accuracy by always choosing `positive` / `negative`. Resampling is what forces it to engage the minority:


- **SMOTE raises `neutral` recall from 0.000 to 0.189** (and `negative` to 0.758) for almost no macro-F1 cost (0.555 → 0.551), because it synthesises minority points rather than discarding majority data.
- **Under-sampling pushes `neutral` recall the highest (0.472)** but, by throwing away majority data, drops accuracy to 0.601 and macro-F1 to 0.505.


This is the textbook lesson of imbalanced learning: **accuracy is the wrong metric** (0.836 with a class never predicted!), and resampling buys genuine minority-class recall at the cost of majority accuracy. SMOTE is the best balance here; if surfacing neutral reviews mattered most, under-sampling would win.

7. Dimensionality reduction (PCA / TruncatedSVD)

The TF-IDF matrix is sparse, so plain PCA (which centres the data) would densify it wastefully. We use **TruncatedSVD** — the PCA-equivalent for sparse text (a.k.a. LSA) — and measure downstream macro-F1 (TF-IDF + Logistic Regression).

Components	Explained variance	Macro-F1
50	28.1 %	0.528
100	37.0 %	0.536
200	48.2 %	0.542
300	56.0 %	0.550
4 630 (full)	100 %	0.555






Interpretation. Compressing the 4 630-dim TF-IDF space to **300 components (6.5 %)** keeps macro-F1 at **0.550 of the 0.555** full-feature score — a $\sim 15\times$ compression for a 0.005 drop, and even 50 components retain 0.528. Dimensionality reduction is therefore an excellent speed/memory trade-off, and is essential before feeding the otherwise-sparse text to a dense learner such as gradient boosting (see §9).

8. Decision-tree pruning

A single Decision Tree on TF-IDF features, traced along its **cost-complexity post-pruning path** (`ccp_alpha`):

<code>ccp_alpha</code>	# nodes	Train acc	Test acc
0.0000 (unpruned)	1 857	0.976	0.759
0.0006	277	0.869	0.783
0.0015	69	0.811	0.786
0.0032 (over-pruned)	29	0.779	0.770




Interpretation. The unpruned tree memorises the training set (train 0.976 vs test 0.759 — a 0.22 generalisation gap). Increasing `ccp_alpha` shrinks it from **1 857 to 69 nodes** while raising test accuracy to **0.786** and shrinking the gap to 0.03; pruning too hard (29 nodes) starts to underfit. Pre-pruning (`max_depth` , `min_samples_leaf`) is also tuned in the main grid (§5). A textbook bias-variance trade-off — a 27× smaller, better-generalising tree.

9. Early stopping (gradient boosting)

Gradient boosting (`GradientBoostingClassifier`) on TF-IDF → TruncatedSVD(100) features, with a generous ceiling of 500 trees and `n_iter_no_change = 10` monitoring a validation slice.

Max trees	Trees actually used	Test accuracy	Test macro-F1
500	88	0.809	0.534



Interpretation. Early stopping halted training after **88 of 500** trees once validation performance plateaued — an **≈ 6× reduction** in training cost with no loss of test quality, and protection against the over-fitting that more trees would bring. (The SVD reduction from §7 is what makes boosting on text efficient.)

10. Discussion: justification of techniques and trade-offs

Decision	Why	Trade-off
Sparse counts (TF-IDF/BoW)	best fit for short lexical review text; TF-IDF+NB won	huge but cheap sparse vocabulary
Keep negation stop-words	preserves sentiment-flipping cues	slightly larger vocabulary
Multinomial vs Gaussian NB by feature type	NB assumptions must match the data	none (handled automatically)
Logistic Regression with L2	strong, robust, regularised baseline	linear decision boundary
TruncatedSVD not PCA	avoids densifying sparse TF-IDF	loses exact feature interpretability
Pre-trained GloVe	tests transfer from general text	mid-table: mean-pooling washes out review-specific cues
Macro-F1 as the metric	the 64 % positive skew makes accuracy meaningless	must be read alongside accuracy
SMOTE inside an imblearn pipeline	rescues zero minority recall, leakage-free	lowers majority accuracy

Decision	Why	Trade-off
Preprocess once + cache vectorisers	whole study in ~8 min	a few hundred MB of disk cache

Headline findings. (1) Simple sparse models win — **TF-IDF + Naive Bayes (0.559 macro-F1)** edges out every embedding/ensemble on short, noisy review text, with the top seven all TF-IDF/BoW. (2) On a 64 %-positive corpus **accuracy (0.84) is a trap**: the untreated baseline never predicts the neutral class at all. (3) **SMOTE is decisive**, lifting `neutral` recall from 0.000 to 0.189 (under-sampling to 0.472) for almost no macro-F1 cost. (4) Boosting and embeddings underperform — AdaBoost stumps and mean-pooled vectors both lose the lexical cues short reviews depend on. (5) Pruning (1 857 → 69 nodes) and early stopping (88/500 trees) both curb over-fitting.

11. Challenges faced

- **Star ratings ≠ clean sentiment labels** — 3★ "neutral" reviews overlap their neighbours and the corpus is extremely positive-skewed (only 4.5 % neutral), making the minority classes (and macro-F1) genuinely hard.
- **Heavy text de-duplication and multilingual noise** — the same short review recurs thousands of times and many reviews are non-English; de-duplicating (to avoid train/test leakage) and dropping empty cleaned text shrinks 10 000 raw rows to 5 923 usable ones.
- **MultinomialNB rejects negative features** → automatic switch to GaussianNB for dense embeddings / reduced features.
- **Run-time of nested tuning** across 20 vectoriser×model combinations → solved with stateless-preprocessing-once + a shared `joblib` vectoriser cache.
- **BERT weights unreachable** behind the network policy → a complete BERT vectoriser is provided but gracefully skipped, with the other four vectorisers (BoW, TF-IDF, Word2Vec, pre-trained GloVe) carrying the comparison.

12. Future work

- Collect **GitHub/JIRA issues** directly once API access is available (the pipeline already accepts any `text,label` CSV), to complement the app reviews.
- Enable the **BERT** path and add a fine-tuned transformer for comparison.
- Try **class-weighting** as a leakage-free alternative to resampling, and cost-sensitive thresholds for the minority class.
- Add **calibrated probabilities** and per-class precision/recall curves for triage use-cases (e.g. auto-escalating angry reports).

13. Reproducibility

```
pip install -r requirements.txt
python main.py                # full study (~8 min on 4 cores), writes results/
```

Every figure and table in this report is regenerated under [results/](#) by that single command, from the seed `RANDOM_STATE = 42`.