

Predicting User Sentiment in Software Issue Reports

A complete, reproducible machine learning pipeline for the
three-class sentiment classification of application reviews

Yann MASSOUAMA · Vénus BAKIKO · Faïçal DIELO

ST2MLE — Machine Learning for IT Engineers

Technical project report

Contents

Abstract	2
1 Introduction	2
1.1 Objectives	2
1.2 Structure of this report	3
2 Problem statement and background	3
2.1 Formal definition	3
2.2 Why the task is non-trivial	3
2.3 Families of text representation	3
3 Dataset and data preparation	4
3.1 Source and labelling convention	4
3.2 De-duplication and cleaning of empty documents	4
3.3 Class distribution	4
4 Methodology	5
4.1 Preprocessing	6
4.2 Vectorisation	6
4.3 Classifiers and their hyper-parameters	7
4.4 Evaluation protocol	7
5 Performance metrics and their meaning	8
6 Results: comparison of all twenty combinations	8
6.1 Analysis of the ranking	10
7 Confusion matrices: every combination in detail	10
7.1 Bag of Words	11
7.2 TF-IDF	12
7.3 In-domain Word2Vec	14
7.4 Pre-trained Word2Vec (Google News)	15
7.5 What the twenty matrices reveal together	17
8 Experiment: handling class imbalance	17
9 Experiment: dimensionality reduction	18
10 Experiment: decision-tree pruning	19
11 Experiment: early stopping	20

12 Discussion and trade-offs	20
13 Conclusion and future work	21

Abstract

This report presents the design, implementation and evaluation of a complete supervised machine learning pipeline for classifying the sentiment of user-written application reviews into three classes: *negative*, *neutral* and *positive*. Working from a real, pre-cleaned corpus of Facebook application reviews, we compare four text-vectorisation strategies (Bag of Words, TF-IDF, an in-domain Word2Vec model and the pre-trained Google News Word2Vec vectors) against five classifiers (Naive Bayes, Decision Tree, Random Forest, AdaBoost and Logistic Regression), for a total of twenty tuned vectoriser–classifier combinations. Each classifier is optimised by grid search under five-fold stratified cross-validation, using the macro-averaged F1 score as the selection criterion because the corpus is strongly imbalanced. Four further experiments study, respectively, class-imbalance handling (SMOTE versus random under-sampling), dimensionality reduction (Truncated SVD), decision-tree pruning (cost-complexity pruning) and early stopping in gradient boosting. The best configuration, TF-IDF with multinomial Naive Bayes, attains a macro-F1 of 0.559 on the held-out test set. Beyond the raw numbers, the study makes a clear methodological point: on an imbalanced dataset, accuracy is deceptive (a model can reach 0.84 accuracy while never predicting the minority class), the macro-F1 and the per-class recall expose this failure, and resampling is what allows the minority class to appear in the predictions, at a measurable cost in majority accuracy.

1 Introduction

Software is now delivered to millions of users who continuously express opinions about it: through application-store reviews, bug reports, feature requests and issue-tracker comments. This feedback is a strategic asset. It tells development teams what users love, what frustrates them, and which defects hurt the product most. The difficulty is one of scale: a popular application receives thousands of reviews per day, far more than any team can read manually. *Sentiment analysis*, the automatic detection of the emotional polarity of a piece of text, is the standard answer to this problem. By labelling each review as negative, neutral or positive, a team can monitor user satisfaction over time, triage the most urgent complaints, and quantify the effect of a release.

This project addresses precisely that task. We build, from raw text to evaluated predictions, a machine learning pipeline that reads the body of an application review and assigns it one of three sentiment classes. The work is deliberately framed as a *comparative study* rather than the pursuit of a single model: the brief for the *ST2MLE — Machine Learning for IT Engineers* course asks not only for a working classifier but for an explicit comparison of several ways of representing text and several classification algorithms, together with cross-validation, hyper-parameter tuning, class-imbalance handling, dimensionality reduction, tree pruning, regularisation and early stopping. Each of these requirements is implemented, measured and discussed.

1.1 Objectives

The objectives of the project are the following.

1. Assemble a reproducible, end-to-end pipeline from raw review text to a final, honestly measured sentiment prediction.
2. Compare four vectorisation strategies and five classifiers under a single, rigorous evaluation protocol, and identify the best combination with evidence.
3. Study, in dedicated experiments, the techniques that the course requires: class-imbalance correction, dimensionality reduction, pruning, regularisation and early stopping.

4. Select and justify metrics that are meaningful on an imbalanced problem, and use them to interpret every result, including the failure modes of the models.

1.2 Structure of this report

Section 2 reviews the problem and the relevant background. Section 3 describes the dataset and its preparation. Section 4 details the full methodology, from preprocessing to the evaluation protocol. Section 5 explains the performance metrics and their meaning in our context. Section 6 presents the main comparison of all twenty combinations. Section 7 reproduces and comments on every confusion matrix individually. Sections 8 to 11 cover the four focused experiments. Section 12 discusses the findings and trade-offs, and Section 13 concludes with perspectives for future work.

2 Problem statement and background

2.1 Formal definition

Let $\mathcal{D} = \{(d_i, y_i)\}_{i=1}^N$ be a corpus of N reviews, where each d_i is a string of natural-language text and each label $y_i \in \{\text{negative}, \text{neutral}, \text{positive}\}$. We seek a function $f : \text{text} \rightarrow \{\text{negative}, \text{neutral}, \text{positive}\}$ that generalises from a training subset to unseen reviews. Because the output space has three mutually exclusive values, this is a *single-label, multi-class* classification problem, as opposed to binary sentiment (positive versus negative) or multi-label tagging. Text cannot be consumed directly by the learning algorithms, so the pipeline first maps each document to a numerical feature vector $\phi(d_i) \in \mathbb{R}^m$ through a *vectoriser*, and the classifier then learns the mapping from $\mathbb{R}^m$ to the label space.

2.2 Why the task is non-trivial

Three properties make this problem genuinely difficult and are worth stating early, because they explain many of the results that follow.

- **Proxy labels.** Our ground-truth labels are derived from star ratings, not from a human annotation of sentiment. A star rating is a convenient but imperfect proxy: a three-star review (mapped to *neutral*) often reads like a faint complaint or a lukewarm endorsement, so the neutral class overlaps its neighbours by construction.
- **Severe class imbalance.** As in most app stores, positive reviews vastly outnumber the rest, and the neutral class is tiny. A classifier can therefore obtain a high accuracy simply by predicting the majority class, which makes accuracy a poor and even misleading objective.
- **Short, informal text.** Reviews are short, contain slang, abbreviations, emoji and spelling mistakes, and are sometimes written in another language. The sentiment of a short review often hinges on one or two words, which has direct consequences for the choice of representation.

2.3 Families of text representation

Turning text into vectors is the central modelling decision, and the literature offers two broad families, both of which we evaluate. *Count-based* representations describe a document by the words it contains. Bag of Words records raw counts; TF-IDF re-weights those counts so that words which are frequent in a document but rare across the corpus (and therefore discriminative) receive more weight. These representations are high-dimensional and sparse, and they preserve the identity of individual words, which matters when a single keyword carries the sentiment. *Embedding-based*

representations map each word to a dense, low-dimensional vector learned so that words used in similar contexts lie close together. Word2Vec is the canonical example. A document can then be represented by the mean of its word vectors. Embeddings capture semantic similarity but, when averaged over a document, they tend to dilute the few strong sentiment words that a short review relies on. This tension between count-based and embedding-based representations is one of the themes of our results.

3 Dataset and data preparation

3.1 Source and labelling convention

We use a real, publicly available dataset of Facebook application reviews exported from Kaggle. Each record contains the review text (the `content` field) and an integer star rating from one to five (the `score` field). We convert the rating into a sentiment label using the standard convention summarised in Table 1. This mapping is deterministic and is the single source of supervision in the project. As noted above, the star rating is a proxy for sentiment, and the three-star to neutral mapping is the principal source of label noise.

Star rating	Assigned sentiment
1 or 2 stars	negative
3 stars	neutral
4 or 5 stars	positive

Table 1: Star-to-sentiment mapping used throughout the project.

3.2 De-duplication and cleaning of empty documents

Application reviews contain a very large number of duplicate short texts; the strings “good”, “nice” or “love it” each appear hundreds of times. We remove duplicate review texts for a subtle but important reason: if an identical string appeared in both the training and the test partition, the model would be evaluated on text it has already seen, and the reported score would be optimistically biased. After de-duplication, and after dropping the reviews that become empty once cleaning removes all of their content (emoji-only reviews, or reviews written entirely in a non-English script), 5 923 unique reviews remain.

3.3 Class distribution

The resulting class distribution, shown in Table 2 and Figure 1, is strongly imbalanced. Positives represent roughly 64% of the corpus, negatives roughly 31%, and the neutral class only about 4.5%. Positives outnumber neutrals by approximately fourteen to one. This imbalance is not an artefact to be removed; it is a realistic property of app-store feedback, and handling it correctly is one of the explicit goals of the study.

Class	From rating	Count	Share
positive	4 to 5 stars	3 801	64.2%
negative	1 to 2 stars	1 857	31.4%
neutral	3 stars	265	4.5%
total		5 923	100%

Table 2: Class distribution after cleaning and de-duplication.




Figure 1: Number of reviews per class on the full corpus. The neutral class is tiny, which drives most of the difficulty analysed later in the report.




Figure 2: Mean number of cleaned tokens per class. The three classes have comparable average lengths, so document length alone is not a discriminative cue and the models must rely on the words themselves.

4 Methodology

The pipeline reads, from left to right: raw text → preprocessing → vectorisation → optional balancing → optional dimensionality reduction → a tuned classifier → evaluation. Two engineering decisions keep the whole study both fast and methodologically sound. First, preprocessing learns nothing from the data (it is *stateless*), so it is applied a single time up front, which is exactly

equivalent to running it inside every cross-validation fold but avoids re-lemmatising the corpus thousands of times. Second, the fitted vectorisers are cached on disk with a shared `joblib.Memory` against a single fixed cross-validation splitter, so the expensive Word2Vec and TF-IDF fits are computed once per fold and reused across every hyper-parameter combination and every classifier. The complete study runs in roughly eight minutes on four cores.

4.1 Preprocessing

Each review is normalised through the classical natural-language sequence *normalise* $\rightarrow$ *tokenise* $\rightarrow$ *remove stop-words* $\rightarrow$ *lemmatise*, with two deliberate, sentiment-aware refinements.

1. **Normalisation.** The text is lower-cased; URLs, HTML tags, punctuation and digits are removed; and runs of whitespace are collapsed. Only alphabetic tokens survive.
2. **Tokenisation.** The cleaned string is split into word tokens.
3. **Stop-word removal with a negation whitelist.** Common, low-information words (“the”, “a”, “is”) are removed, *except* negation cues such as “not”, “no”, “never”, “without” and “cannot”. This exception is important: the phrases “working” and “not working” carry opposite sentiment, and discarding “not” as an ordinary stop-word would destroy that signal.
4. **Lemmatisation.** A light two-pass WordNet lemmatisation maps each token first to its verb lemma and then to its noun lemma, so that “crashing” becomes “crash” and “issues” becomes “issue”. This shrinks the vocabulary and groups inflected forms of the same word without the computational cost of a full part-of-speech tagger.

As a concrete example, the raw review “Critical issue, the app keeps crashing!!!” is transformed into “critical issue app keep crash”.

4.2 Vectorisation

We compare the four representations named in the brief, summarised in Table 3. Bag of Words and TF-IDF both produce high-dimensional, sparse and non-negative vectors; they differ only in the weighting. For TF-IDF we enable sublinear term-frequency scaling, a minimum document frequency of two (so that terms appearing in a single document are dropped as noise), a vocabulary capped at the 5 000 most frequent terms, and an n-gram range of (1, 2) so that informative word pairs such as “not good” are captured alongside single words. The in-domain Word2Vec model is trained on our own reviews (100-dimensional vectors, context window of five, thirty epochs); a document is represented by the mean of its word vectors. The pre-trained variant reuses the publicly available Google News Word2Vec vectors (300-dimensional, trained on roughly one hundred billion words), which is the exact *pre-trained Word2Vec* the brief refers to. Keeping both an in-domain and a pre-trained embedding lets us contrast learning representations on our own, limited data against reusing representations learned on a very large external corpus, a form of transfer learning.

Method	Underlying idea	Output
Bag of Words	raw count of each word	sparse, non-negative
TF-IDF	counts re-weighted by inverse document frequency	sparse, non-negative
Word2Vec (in-domain)	mean of embeddings trained on our reviews	dense
Word2Vec (pre-trained)	mean of Google News 300d vectors	dense

Table 3: The four vectorisers. Sparse, non-negative outputs feed multinomial Naive Bayes; dense outputs use the Gaussian variant, selected automatically.

A BERT vectoriser is also implemented and documented in the code base. It requires the optional deep-learning stack and network access to the model hub, which is not reachable in our execution environment, so it is skipped automatically and provided through a dedicated script that runs wherever the hub is available, for example on Google Colab.

4.3 Classifiers and their hyper-parameters

We compare five classifiers, chosen so that together they cover the techniques the brief requires (pruning, ensembles, boosting and regularisation). For each one, the hyper-parameters listed in Table 4 are tuned by grid search; the values were chosen to be broad enough to demonstrate meaningful tuning yet compact enough to remain fast under five-fold cross-validation.

Classifier	Tuned hyper-parameter	Values searched
Naive Bayes (multinomial)	<code>alpha</code> (Laplace smoothing)	0.1, 0.5, 1.0
Naive Bayes (Gaussian)	<code>var_smoothing</code>	10^{-9} , 10^{-8} , 10^{-7}
Decision Tree	<code>max_depth</code> (pre-pruning)	None, 10, 20, 30
	<code>min_samples_leaf</code> (pre-pruning)	1, 5, 10
	<code>ccp_alpha</code> (post-pruning)	0, 10^{-3} , 10^{-2}
Random Forest	<code>n_estimators</code>	200
	<code>max_depth</code>	None, 20
	<code>min_samples_leaf</code>	1, 2
	<code>max_features</code>	<code>sqrt</code>
AdaBoost	<code>n_estimators</code>	100, 200
	<code>learning_rate</code>	0.5, 1.0
Logistic Regression	<code>C</code> (inverse L2 strength)	0.1, 1.0, 10.0

Table 4: Hyper-parameter grids per classifier. The Naive Bayes variant is selected automatically: multinomial for sparse count features, Gaussian for dense embeddings.

The intuition behind each model is as follows. **Naive Bayes** estimates, for each class, the probability of each word, and predicts the class that maximises the product of the prior and the word likelihoods; it assumes the words are conditionally independent, which is false in practice but works very well and is extremely fast in high dimension. The Laplace smoothing parameter `alpha` prevents a single unseen word from forcing a zero probability. **The Decision Tree** is interpretable but tends to over-fit, which is exactly why both its pre-pruning parameters (`max_depth`, `min_samples_leaf`) and its post-pruning parameter (`ccp_alpha`) are placed in the grid. **The Random Forest** averages many de-correlated trees and is a strong low-variance ensemble. **AdaBoost** combines many very shallow trees, re-weighting the hard cases at each round. **Logistic Regression** draws a regularised linear boundary; tuning its inverse-strength parameter `C` is our explicit demonstration of L2 regularisation, where a small `C` regularises strongly (a simpler model, less over-fitting) and a large `C` lets the model fit the data more closely.

4.4 Evaluation protocol

The data is split into 80% training and 20% test in a stratified manner, so that the three class proportions are preserved in both partitions. On the training set we run five-fold stratified cross-validation and tune every classifier with a grid search that maximises the macro-averaged F1 score. Crucially, the held-out test set is touched only once, at the very end, to report the final numbers,

and any resampling is performed inside the pipeline so that it acts on the training folds alone and never on the validation or test data. Every stochastic component is seeded, so a clean run reproduces the same results exactly. The five-fold cross-validation, illustrated conceptually in the rotation where each of the five folds serves once as the validation set while the other four train the model, yields a stable and honest estimate of generalisation performance and prevents the hyper-parameter choice from over-fitting a single lucky split.

5 Performance metrics and their meaning

Because the problem has three classes, each metric is defined per class (one class against the rest) and then averaged. For a given class, a *true positive* is a review that truly belongs to the class and is predicted as such; a *false positive* is a review predicted as the class but truly belonging to another; and a *false negative* is a true member of the class that the model assigned elsewhere.

- **Accuracy** is the proportion of correctly classified reviews over the whole test set. On an imbalanced problem it is misleading, because a model can score high by always predicting the dominant class.
- **Precision** of a class, $TP/(TP + FP)$, answers: when the model predicts this class, how often is it right? It measures the reliability of the predictions for that class.
- **Recall** of a class, $TP/(TP + FN)$, answers: among the reviews that truly belong to this class, how many does the model recover? It measures coverage.
- **F1** is the harmonic mean of precision and recall, $2PR/(P + R)$, so it is low whenever either of the two is low; it cannot be high unless both are.
- **Macro-F1** is the unweighted mean of the three per-class F1 scores. It treats every class equally regardless of its size, so neglecting the rare class is penalised heavily.

This is precisely why we tune on the macro-F1 rather than on accuracy. On a corpus that is 64% positive, accuracy rewards a model that ignores the minority classes, whereas the macro-F1 and the per-class recall expose that behaviour immediately. The *macro* average (each class counts equally) must be distinguished from the *micro* average (each review counts equally, which on this corpus is dominated by the positive class and behaves like accuracy). Throughout the report, the macro-F1 and the per-class recall are the headline numbers; accuracy is reported for completeness but is never the basis of a decision.

6 Results: comparison of all twenty combinations

Table 5 reports the test-set results for all twenty vectoriser–classifier combinations, sorted by macro-F1. The column **CV F1** is the mean macro-F1 across the five cross-validation folds on the training set (the criterion used for tuning), while the remaining columns are computed once on the held-out test set. Figures 3 and 4 display the same information graphically.

#	Vectoriser	Classifier	CV F1	Acc.	Prec.	Rec.	F1 macro
1	TF-IDF	Naive Bayes	0.550	0.839	0.550	0.569	0.559
2	TF-IDF	Logistic Regression	0.557	0.830	0.547	0.563	0.555
3	TF-IDF	Random Forest	0.541	0.824	0.540	0.558	0.548
4	BoW	Naive Bayes	0.569	0.817	0.542	0.554	0.548
5	BoW	Logistic Regression	0.561	0.806	0.550	0.545	0.546
6	Word2Vec pre-trained	Random Forest	0.535	0.821	0.542	0.545	0.542
7	BoW	Decision Tree	0.533	0.768	0.547	0.533	0.539
8	BoW	Random Forest	0.542	0.810	0.534	0.542	0.537
9	Word2Vec	Logistic Regression	0.532	0.810	0.527	0.547	0.537
10	Word2Vec pre-trained	Logistic Regression	0.558	0.802	0.525	0.544	0.535
11	Word2Vec	Random Forest	0.539	0.801	0.520	0.542	0.530
12	TF-IDF	Decision Tree	0.527	0.784	0.578	0.530	0.528
13	Word2Vec	Naive Bayes	0.517	0.685	0.542	0.537	0.518
14	Word2Vec	AdaBoost	0.518	0.780	0.503	0.531	0.516
15	Word2Vec pre-trained	AdaBoost	0.527	0.783	0.506	0.527	0.516
16	Word2Vec	Decision Tree	0.533	0.763	0.496	0.514	0.505
17	Word2Vec pre-trained	Decision Tree	0.506	0.746	0.483	0.504	0.493
18	TF-IDF	AdaBoost	0.459	0.757	0.512	0.471	0.472
19	BoW	AdaBoost	0.454	0.753	0.516	0.463	0.465
20	Word2Vec pre-trained	Naive Bayes	0.444	0.630	0.465	0.479	0.455

Table 5: Full comparison on the held-out test set, sorted by macro-F1. The best configuration is TF-IDF with Naive Bayes.




Figure 3: Macro-F1 of every classifier, grouped by vectoriser.




Figure 4: Accuracy of every classifier, grouped by vectoriser. Note how the accuracy ranking differs from the macro-F1 ranking, a direct consequence of the class imbalance.

6.1 Analysis of the ranking

Several lessons stand out. The best combination is TF-IDF with Naive Bayes at 0.559 macro-F1, and the five best entries are all built on sparse count features (TF-IDF or Bag of Words). On short and noisy review text, counting the right words and weighting them by rarity beats averaging dense embeddings, because the mean of word vectors tends to wash out the few strong sentiment words on which a short review depends. AdaBoost sits at the very bottom with both count representations, because its shallow decision stumps can only test one feature at a time out of several thousand, which is a poor match for high-dimensional sparse text. The pre-trained Word2Vec behaves reasonably with the Random Forest, where it reaches 0.542 and ranks sixth, yet it is the single worst entry when paired with Gaussian Naive Bayes at 0.455; this is a useful reminder that a pre-trained representation is not automatically superior to one learned on the data at hand, and that the classifier has to match the geometry of the features. Finally, the scores cluster tightly between 0.46 and 0.56. This narrow spread is the honest signature of a real and difficult problem: no representation is dramatically better than the others, and the ceiling is set by the ambiguity of the labels rather than by the choice of algorithm.

7 Confusion matrices: every combination in detail

This section reproduces and comments on the confusion matrix of every one of the twenty combinations. In each matrix the rows are the true classes and the columns are the predicted classes, in the fixed order negative, neutral, positive, so the diagonal holds the correct predictions. Reading these matrices, rather than a single aggregate score, is what reveals *how* each model succeeds or fails, and in particular how it treats the tiny neutral class. For reference, the held-out test set contains 372 negative, 53 neutral and 760 positive reviews, for a total of 1185 reviews; the training set on which the models are fitted contains 1485, 212 and 3041 reviews of each class respectively. In every matrix below, the rows sum to these true class counts, which is a quick sanity check that the matrix is internally consistent.

A recurring pattern, visible across almost all twenty matrices, is that the neutral column is empty or nearly empty: only five of the twenty models ever assign the neutral label correctly even once,

and only two do so more than a handful of times. This is the *accuracy paradox* in action, and the per-vectoriser tables below make it explicit by reporting the F1 of each class separately.

7.1 Bag of Words

Table 6 summarises the five Bag of Words models. The neutral F1 is zero or close to zero everywhere, and the macro-F1 is driven almost entirely by the negative and positive F1 scores.

Classifier	Acc.	F1 neg.	F1 neu.	F1 pos.	Macro F1
Naive Bayes	0.817	0.762	0.000	0.882	0.548
Decision Tree	0.768	0.680	0.089	0.847	0.539
Random Forest	0.810	0.737	0.000	0.875	0.537
AdaBoost	0.753	0.554	0.000	0.840	0.465
Logistic Regression	0.806	0.735	0.026	0.877	0.546

Table 6: Per-class F1 for the five Bag of Words models on the test set.




Figure 5: Confusion matrices of the five Bag of Words models. Rows are true classes, columns predicted classes, ordered negative, neutral, positive.

BoW + Naive Bayes. The model correctly labels 283 of 372 negatives (recall 0.761) and 685 of 760 positives (recall 0.901), but not one of the 53 neutrals. Of the 20 reviews it does send to the neutral column, 13 are truly negative and 7 truly positive, so its rare neutral guesses are always wrong, giving a neutral precision and recall of zero. The result is a respectable macro-F1 of 0.548 built entirely on the two large classes.

BoW + Decision Tree. This is one of the few models that recovers any neutral review at all: 4 of 53 (recall 0.075, F1 0.089). It pays a clear price for this: 113 negatives leak into the positive

column, dropping the negative recall to 0.651, and it predicts neutral 37 times for only 4 correct hits, a neutral precision of just 0.108. The tree is the noisiest BoW model but, paradoxically, the most willing to bet on the minority class.

BoW + Random Forest. Averaging many trees makes the forest even more reluctant than a single tree to predict neutral: the neutral column holds 9 predictions, none of them correct. Negative and positive recalls are solid (0.712 and 0.914), and the macro-F1 of 0.537 is typical of the strong-but-neutral-blind behaviour seen throughout.

BoW + AdaBoost. This matrix exposes the weakness of boosting shallow stumps on sparse text. Not only is the neutral column entirely empty, but the negative class also collapses: only 157 of 372 negatives are recovered (recall 0.422) while 215 are flipped to positive. The positive recall is inflated to 0.967, and the macro-F1 sinks to 0.465, the worst of the Bag of Words group. Because each stump can test only one word at a time, the ensemble defaults to the dominant positive class.

BoW + Logistic Regression. The linear model behaves like Naive Bayes but is a touch more balanced, recovering 264 negatives (recall 0.710) and 690 positives (recall 0.908), with a single correct neutral (recall 0.019). Its macro-F1 of 0.546 makes it the second-best Bag of Words model, just behind Naive Bayes.

7.2 TF-IDF

TF-IDF produces the best models of the entire study. Table 7 shows that Naive Bayes and Logistic Regression lead, while AdaBoost again trails badly.

Classifier	Acc.	F1 neg.	F1 neu.	F1 pos.	Macro F1
Naive Bayes	0.839	0.784	0.000	0.893	0.559
Decision Tree	0.784	0.700	0.034	0.850	0.528
Random Forest	0.824	0.763	0.000	0.881	0.548
AdaBoost	0.757	0.576	0.000	0.842	0.472
Logistic Regression	0.830	0.776	0.000	0.889	0.555

Table 7: Per-class F1 for the five TF-IDF models on the test set.




Figure 6: Confusion matrices of the five TF-IDF models.

TF-IDF + Naive Bayes (best overall). This is the canonical accuracy-paradox matrix and the best model of the study. The neutral column is exactly zero: the 53 neutral reviews are split between negative (22) and positive (31), and the model never predicts neutral. Yet it is excellent on the two large classes, with 291 negatives (recall 0.782, precision 0.786) and 703 positives (recall 0.925, precision 0.863). The accuracy reaches 0.839 and the weighted F1 about 0.82, while the macro-F1 is only 0.559 precisely because the neutral F1 is zero. It is the clearest illustration of why we report the macro-F1.

TF-IDF + Decision Tree. The tree recovers a single neutral (recall 0.019) at a precision of 0.20, and otherwise loses ground on the negative class, with 116 negatives leaking to positive (recall 0.683). Its macro-F1 of 0.528 is the lowest among the TF-IDF models that are not AdaBoost, a reminder that an unpruned-style tree over-fits the training data and generalises less well than the linear and probabilistic models.

TF-IDF + Random Forest. A clean, strong matrix: 282 negatives (recall 0.758) and 695 positives (recall 0.914) recovered, neutral column empty. The forest matches Naive Bayes on the macro-F1 (0.548) and is the most reliable ensemble here, though it shares the universal blindness to neutral.

TF-IDF + AdaBoost. The boosting pathology recurs identically to the Bag of Words case: 203 of 372 negatives are flipped to positive (negative recall 0.454), the neutral column is empty, and the positive recall is inflated to 0.958. The macro-F1 of 0.472 confirms that AdaBoost on shallow stumps is structurally ill-suited to wide sparse inputs, regardless of whether the features are raw counts or TF-IDF weights.

TF-IDF + Logistic Regression. Almost a mirror of the best Naive Bayes model: 289 negatives (recall 0.777) and 694 positives (recall 0.913) recovered, neutral essentially absent. Its macro-F1

of 0.555 ranks second overall and shows that a regularised linear boundary and a generative word model arrive at very similar, strong solutions on this representation.

7.3 In-domain Word2Vec

Averaging embeddings learned on our own corpus shifts the behaviour noticeably, as Table 8 shows. Most strikingly, this is the family where one model finally engages the neutral class in a meaningful way.

Classifier	Acc.	F1 neg.	F1 neu.	F1 pos.	Macro F1
Naive Bayes	0.685	0.628	0.110	0.816	0.518
Decision Tree	0.763	0.674	0.000	0.840	0.505
Random Forest	0.801	0.728	0.000	0.863	0.530
AdaBoost	0.780	0.702	0.000	0.846	0.516
Logistic Regression	0.810	0.739	0.000	0.871	0.537

Table 8: Per-class F1 for the five in-domain Word2Vec models on the test set.




Figure 7: Confusion matrices of the five in-domain Word2Vec models.

Word2Vec + Naive Bayes (the neutral-aware outlier). This is the most unusual matrix in the study and deserves close attention. The Gaussian Naive Bayes on dense embeddings is the only model that genuinely engages the neutral class, correctly labelling 14 of 53 neutrals, a recall of 0.264 that dwarfs every other model. It achieves this because a generative per-class density on continuous features gives the small class a real region of its own, rather than letting a dominant prior crush it. The price is heavy, however: it predicts neutral 202 times in total, so the neutral precision is only 0.069, and it misreads 95 positives as neutral, pulling the positive recall down

to 0.766 and the overall accuracy down to 0.685, the lowest of the in-domain group. It trades precision and majority accuracy for minority coverage.

Word2Vec + Decision Tree. Back to a blind-to-neutral pattern: zero neutrals recovered, 109 negatives lost to the positive column. With a macro-F1 of 0.505 it is the weakest in-domain model, confirming that a single tree on a 100-dimensional dense space over-fits without adding minority sensitivity.

Word2Vec + Random Forest. The strongest in-domain model at macro-F1 0.530: 274 negatives (recall 0.737) and 675 positives (recall 0.888) recovered, neutral absent. The ensemble stabilises the noisy single-tree behaviour but, as everywhere, does not rescue the minority class.

Word2Vec + AdaBoost. Interestingly, on dense embeddings AdaBoost keeps a healthy negative recall (0.737), unlike its collapse on sparse counts, because each stump now splits on a continuous embedding dimension rather than a single rare word. The neutral column is still empty, and the macro-F1 is 0.516.

Word2Vec + Logistic Regression. A balanced, neutral-blind matrix with negative and positive recalls of 0.739 and 0.901, for a macro-F1 of 0.537, essentially tied with the Random Forest at the top of the in-domain group.

7.4 Pre-trained Word2Vec (Google News)

Reusing the Google News vectors brings external knowledge but, as Table 9 shows, it does not lift the minority class and produces both the sixth-best and the single worst model of the study, depending on the classifier.

Classifier	Acc.	F1 neg.	F1 neu.	F1 pos.	Macro F1
Naive Bayes	0.630	0.627	0.044	0.695	0.455
Decision Tree	0.746	0.658	0.000	0.821	0.493
Random Forest	0.821	0.741	0.000	0.884	0.542
AdaBoost	0.783	0.697	0.000	0.851	0.516
Logistic Regression	0.802	0.736	0.000	0.867	0.535

Table 9: Per-class F1 for the five pre-trained Word2Vec models on the test set.




Figure 8: Confusion matrices of the five pre-trained Word2Vec (Google News) models.

Pre-trained Word2Vec + Naive Bayes (worst overall). With a macro-F1 of 0.455 this is the weakest model in the study, and its matrix explains why. The Gaussian model over-predicts the negative class: it recovers 293 negatives (recall 0.788) but misclassifies 240 true positives as negative, collapsing the positive recall to 0.593 and the accuracy to 0.630. It labels three neutrals correctly (recall 0.057). The general purpose Google vectors, optimised for word analogy rather than sentiment, do not separate the sentiment classes cleanly enough for a generative density model.

Pre-trained Word2Vec + Decision Tree. Zero neutrals, and now positives leak the other way, with 128 positives sent to the negative column. The macro-F1 of 0.493 is the second-lowest overall, confirming that a single tree handles the dense pre-trained space poorly.

Pre-trained Word2Vec + Random Forest (best pre-trained model). This is the strongest pre-trained configuration and ranks sixth overall at macro-F1 0.542. It recovers 259 negatives (recall 0.696) and 714 positives (recall 0.939, the highest positive recall among the pre-trained models), with an empty neutral column. The pairing of an ensemble with transfer-learned features works reasonably well, which is why it is the only embedding model to break into the top six.

Pre-trained Word2Vec + AdaBoost. As with the in-domain embeddings, AdaBoost keeps a reasonable negative recall (0.704) on the dense features, the neutral column stays empty, and the macro-F1 is 0.516, essentially identical to its in-domain counterpart.

Pre-trained Word2Vec + Logistic Regression. A balanced, neutral-blind matrix with negative and positive recalls of 0.750 and 0.883 and a macro-F1 of 0.535, the second best pre-trained model and very close to its in-domain Logistic Regression twin.

7.5 What the twenty matrices reveal together

Reading all twenty matrices side by side yields three robust conclusions that no single aggregate number conveys. First, the neutral class is almost universally sacrificed: fifteen of the twenty models never predict it correctly even once, and the macro-F1 is therefore governed by how well a model balances the negative and positive classes. Second, the only models that engage the neutral class at all are those built on dense embeddings with a generative Gaussian Naive Bayes, and they do so at a steep cost in precision and majority accuracy, which previews exactly why the dedicated imbalance experiment in Section 8 is necessary. Third, AdaBoost on sparse count features fails in a characteristic way, collapsing the negative class into the positive one, whereas on dense embeddings it recovers a normal negative recall; the failure is therefore an interaction between the algorithm and the feature geometry, not a property of the algorithm alone. These observations, invisible in a table of macro-F1 scores, are precisely what the confusion matrices are for.

8 Experiment: handling class imbalance

To address the imbalance directly, we compare three strategies on a fixed pipeline that uses TF-IDF features and a Logistic Regression classifier, so that the only thing that changes is the resampling step. The sampler is placed inside an `imblearn` pipeline, which guarantees that it acts on the training folds alone and is bypassed at prediction time, eliminating any leakage of synthetic samples into the evaluation.

Strategy	Accuracy	Macro F1	Rec. neg.	Rec. neu.	Rec. pos.
none (baseline)	0.836	0.555	0.745	0.000	0.939
SMOTE	0.727	0.551	0.758	0.189	0.749
under-sampling	0.601	0.505	0.637	0.472	0.592

Table 10: Effect of resampling. The neutral recall is the column to watch: it rises from zero to 0.189 with SMOTE and to 0.472 with under-sampling.




Figure 9: Macro-F1 for each resampling strategy.

This is among the most instructive results of the study. Without any treatment, the baseline reaches 0.836 accuracy while its neutral recall is exactly zero: it never recognises a single neutral review. SMOTE synthesises new neutral examples by interpolating between neighbouring minority points (with five neighbours), and it lifts the neutral recall from zero to 0.189 at almost no cost in

macro-F1, because it adds minority points rather than discarding majority data. Random under-sampling pushes the neutral recall much further, to 0.472, but it does so by throwing away the majority examples it needs, which collapses the accuracy to 0.601 and the macro-F1 to 0.505. The honest conclusion is that resampling does not create free performance; it *redistributes* errors, trading majority accuracy for genuine minority recall. SMOTE offers the best overall balance here, whereas under-sampling would be preferable if the explicit priority were to surface neutral reviews even at the expense of overall accuracy. The choice of strategy is therefore a decision driven by the application’s goal, not a purely technical optimisation.

9 Experiment: dimensionality reduction

The TF-IDF matrix has about 4 600 columns, one per vocabulary term. We reduce this dimensionality with Truncated SVD, which is the variant of principal component analysis suited to sparse text. Ordinary PCA centres the data by subtracting the column means; on a sparse matrix this turns every stored zero into a non-zero value and densifies the matrix, which is wasteful and memory-hungry. Truncated SVD performs the same singular value decomposition without centring, so the matrix stays sparse; in the natural language processing literature this is known as Latent Semantic Analysis. We then measure the downstream macro-F1 of a Logistic Regression trained on the reduced features.

Components	Cumulative explained variance	Macro F1
50	28.1%	0.528
100	37.0%	0.536
200	48.2%	0.542
300	56.0%	0.550
4 630 (full)	100%	0.555

Table 11: Downstream performance against the number of retained SVD components.




Figure 10: Truncated SVD: information retained (left) and downstream performance (right).

The explained-variance curve rises steeply at first and then flattens, the classic concave shape: the first components capture most of the structure and each additional component contributes less. The performance curve is almost flat. Keeping only 300 components, about six percent of the features, preserves a macro-F1 of 0.550 against 0.555 for the full representation, a loss of only 0.005, and even 50 components retain 0.528. Dimensionality reduction is therefore an excellent

trade-off between speed and accuracy on this problem, and it becomes a necessity, not a luxury, before feeding the otherwise sparse text to a dense model such as gradient boosting, which is the subject of the early-stopping experiment.

10 Experiment: decision-tree pruning

A single decision tree trained on TF-IDF features illustrates pruning very clearly. We follow the cost-complexity pruning path, in which a single parameter `ccp_alpha` controls how aggressively the least useful branches are removed: a larger value prunes more and yields a smaller tree. Table 12 reports representative points along the path, and Figure 11 shows the full train and test accuracy curves.

<code>ccp_alpha</code>	Number of nodes	Train accuracy	Test accuracy
0.0000 (unpruned)	1 857	0.976	0.759
0.0006	277	0.869	0.783
0.0015	69	0.811	0.786
0.0032 (over-pruned)	29	0.779	0.770

Table 12: Train and test accuracy along the cost-complexity pruning path.




Figure 11: Train and test accuracy as the pruning strength `ccp_alpha` increases. The two curves form the characteristic scissor shape of the bias–variance trade-off.

The unpruned tree memorises the training set: it reaches 0.976 accuracy in training but only 0.759 in test, and the large gap of about 0.22 is a textbook signature of over-fitting, in a tree of 1 857 nodes that is far too large to interpret. Pruning reduces the tree from 1 857 to just 69 nodes; the training accuracy falls to 0.811 but the *test* accuracy rises to 0.786, and the train–test gap shrinks from about 0.22 to 0.03. We have traded memorisation for generalisation, which is exactly the goal, and the falling training accuracy is here a sign of success, not of failure. Pruning too aggressively, down to 29 nodes, begins to under-fit and the test accuracy slips back to 0.770. The best tree, at 69 nodes, sits at the bottom of the bias–variance valley: complex enough to capture the signal, simple enough to generalise, and, as a bonus, small enough to be read and understood.

11 Experiment: early stopping

For the gradient boosting model, which builds an additive ensemble by adding one tree at a time, we allow a generous ceiling of 500 trees, reduce the TF-IDF features to 100 SVD components so that the dense booster trains efficiently, and monitor a held-out validation slice, stopping as soon as the validation performance stops improving for ten consecutive rounds. Training halts automatically after only 88 trees out of the 500 allowed, roughly six times less computation, with no loss of test quality and a clear protection against the over-fitting that the additional trees would eventually introduce.




Figure 12: Test accuracy as a function of the number of boosting iterations. The accuracy climbs quickly and then plateaus; the model stops well before the ceiling, at 88 trees, once additional trees no longer improve the validation score.

The shape of the curve is the message: the accuracy climbs quickly and then reaches a plateau, after which more trees add cost without benefit. Early stopping detects this plateau and halts training, so the final number of trees is selected automatically by the data rather than fixed arbitrarily by the practitioner. The four ablation experiments (imbalance, dimensionality reduction, pruning and early stopping) thus share a single underlying theme: each one is, in its own way, a search for the right amount of model complexity, enough to learn the signal but not so much as to memorise the noise.

12 Discussion and trade-offs

Taken together, the experiments paint a coherent picture. Simple sparse models win on short and noisy review text, and TF-IDF with Naive Bayes is the sweet spot of the whole study. Accuracy is the wrong headline metric on this dataset, since a model can reach 0.84 accuracy while never predicting the neutral class at all, so the macro-F1 and the per-class recall are the numbers we report and discuss. Resampling, and SMOTE in particular, is what allows the minority class to exist in the predictions, at the cost of some majority accuracy; dimensionality reduction buys efficiency for a negligible loss in performance; and pruning and early stopping both curb over-fitting in a measurable and visible way. The embeddings, whether trained in-domain or pre-trained on Google News, sit in the middle of the ranking, because averaging word vectors dilutes the few strong sentiment words that short reviews live on.

It is worth being explicit about the advantages and limitations of each representation. Bag of Words and TF-IDF are fast, interpretable (one can read off the influential words) and strong on this task, but they ignore word order beyond the bigrams we include and they produce very high-

dimensional spaces. The in-domain Word2Vec captures some semantic similarity and is compact, but it is starved of data on a corpus of a few thousand reviews. The pre-trained Word2Vec brings external knowledge through transfer learning, yet its general-purpose vectors are not tuned to the sentiment task and, once averaged, lose the local cues that matter. Among classifiers, Naive Bayes is fast and surprisingly strong despite its independence assumption; Logistic Regression is a robust, well-calibrated linear baseline; the Random Forest is reliable but heavier; the single Decision Tree is interpretable but needs pruning to generalise; and AdaBoost, with its one-feature stumps, is simply ill-suited to wide sparse inputs.

13 Conclusion and future work

We have delivered a complete and reproducible sentiment-classification pipeline on a real dataset of Facebook application reviews. We compared four vectorisers and five classifiers, for twenty tuned combinations evaluated under a single rigorous protocol, and we studied class-imbalance handling, dimensionality reduction, pruning and early stopping in dedicated experiments. The best model is TF-IDF with multinomial Naive Bayes, at a macro-F1 of 0.559. Beyond the numbers, the project makes a clear methodological point that generalises well beyond this dataset: on an imbalanced problem the accuracy can be deceptive, the macro-F1 and the per-class recall tell the truth, and resampling is what gives the minority class a voice.

The main difficulties were honest properties of real data rather than software defects. The star-based labels are only a proxy for sentiment; the neutral class is both tiny and intrinsically ambiguous; and the text is short, informal and occasionally written in another language. Heavy de-duplication was necessary to avoid leakage between the training and the test sets. Several avenues would extend this work. First, enabling the BERT vectoriser, already implemented and runnable through a dedicated script wherever the model hub is reachable, would test whether contextual embeddings overcome the dilution that averaged Word2Vec suffers from. Second, class weighting and adjusted decision thresholds are natural alternatives to resampling and could lift minority recall with a different accuracy trade-off. Third, collecting genuine GitHub or JIRA issue reports, rather than app-store reviews, would broaden the pipeline towards the full range of software issue reports named in the project title. Finally, a small campaign of human annotation on the three-star reviews would quantify, and perhaps reduce, the label noise that currently caps the achievable performance.